

project12

February 13, 2026

0.1 Customer Churn Prediction Project

0.1.1 (Machine Learning Project (Phase 3))

1 Project Overview

This project focuses on predicting customer churn using machine learning models. Predicting churn is a critical business objective because retaining existing customers is far more cost-effective than acquiring new ones. By identifying customers at risk of leaving, the company can deploy proactive retention strategies such as personalized discounts, loyalty programs, or targeted engagement campaigns to preserve revenue and maintain market share. This analysis evaluates four models: - Logistic Regression (baseline linear model) - Decision Tree Classifier (rule-based model) - Random Forest (ensemble of decision trees) - Gradient Boosting (boosted ensemble model)

The goal is to determine which model best identifies customers likely to churn while balancing interpretability and predictive performance.

Load all necessary libraries

```
[1]: # Importing necessary libraries
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.preprocessing import LabelBinarizer, MinMaxScaler, LabelEncoder, StandardScaler
from sklearn.linear_model import Ridge, LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error, r2_score, classification_report, recall_score, precision_score, f1_score, accuracy_score, confusion_matrix, ConfusionMatrixDisplay, roc_curve, auc, RocCurveDisplay
from sklearn.feature_selection import RFE
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
import warnings
warnings.filterwarnings('ignore')
```

Load Churn dataset

```
[2]: #Load data set and display the first five datasets through .head() function
df = pd.read_csv('Churn.csv')
df.head()
```

```
[2]: state account length area code phone number international plan \
0 KS 128 415 382-4657 no
1 OH 107 415 371-7191 no
2 NJ 137 415 358-1921 no
3 OH 84 408 375-9999 yes
4 OK 75 415 330-6626 yes

voice mail plan number vmail messages total day minutes total day calls \
0 yes 25 265.1 110
1 yes 26 161.6 123
2 no 0 243.4 114
3 no 0 299.4 71
4 no 0 166.7 113

total day charge ... total eve calls total eve charge \
0 45.07 ... 99 16.78
1 27.47 ... 103 16.62
2 41.38 ... 110 10.30
3 50.90 ... 88 5.26
4 28.34 ... 122 12.61

total night minutes total night calls total night charge \
0 244.7 91 11.01
1 254.4 103 11.45
2 162.6 104 7.32
3 196.9 89 8.86
4 186.9 121 8.41

total intl minutes total intl calls total intl charge \
0 10.0 3 2.70
1 13.7 3 3.70
2 12.2 5 3.29
3 6.6 7 1.78
4 10.1 3 2.73

customer service calls churn
0 1 False
1 1 False
2 0 False
3 2 False
4 3 False
```

[5 rows x 21 columns]

Understanding the data

```
[3]: #Use of .info()
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3333 entries, 0 to 3332
Data columns (total 21 columns):
 #   Column           Non-Null Count  Dtype  
 --- 
 0   state            3333 non-null    object  
 1   account length   3333 non-null    int64  
 2   area code         3333 non-null    int64  
 3   phone number     3333 non-null    object  
 4   international plan 3333 non-null    object  
 5   voice mail plan  3333 non-null    object  
 6   number vmail messages 3333 non-null    int64  
 7   total day minutes 3333 non-null    float64 
 8   total day calls   3333 non-null    int64  
 9   total day charge   3333 non-null    float64 
 10  total eve minutes 3333 non-null    float64 
 11  total eve calls   3333 non-null    int64  
 12  total eve charge   3333 non-null    float64 
 13  total night minutes 3333 non-null    float64 
 14  total night calls  3333 non-null    int64  
 15  total night charge  3333 non-null    float64 
 16  total intl minutes 3333 non-null    float64 
 17  total intl calls   3333 non-null    int64  
 18  total intl charge   3333 non-null    float64 
 19  customer service calls 3333 non-null    int64  
 20  churn             3333 non-null    bool    
dtypes: bool(1), float64(8), int64(8), object(4)
memory usage: 524.2+ KB
```

Run the .column to confirm my data columns


```
[4]: df.columns
```

```
[4]: Index(['state', 'account length', 'area code', 'phone number',
       'international plan', 'voice mail plan', 'number vmail messages',
       'total day minutes', 'total day calls', 'total day charge',
       'total eve minutes', 'total eve calls', 'total eve charge',
       'total night minutes', 'total night calls', 'total night charge',
       'total intl minutes', 'total intl calls', 'total intl charge',
       'customer service calls', 'churn'],
      dtype='object')
```

EDA Chart -> Helps to understand data set

```
[5]: # Churn vs not churn bar chart


df['churn'].value_counts().plot(kind='bar', color='red')
plt.title("Churn Distribution")
plt.xlabel("Churn (0=No, 1=Yes)")
plt.ylabel("Count")
plt.show()
```



Explanation - The chart above confirms imbalanced data hence using accuracy score alone will not help

```
[6]: #Distribution of total day minutes


plt.Figure(figsize=(10,8))
plt.hist(df['total day minutes'], bins=25)
plt.title("Distribution of Total Day Minutes")
plt.xlabel("Minutes")
plt.ylabel("Frequency")
plt.show()
```



Explanation - The chart above helps to identify skewness, outlier and how data is spread

```
[7]: #Box plot showing total day minutes by churn
plt.figure(figsize=(20,16))
df.boxplot(column='total day minutes', by='churn')
plt.title("Total Day Minutes by Churn")
plt.suptitle("")
```

<Figure size 1440x1152 with 0 Axes>




Explanation - The chart above helps to determine if churn customers use more minutes

```
[8]: #Comparing churn rate with customer service calls

grouped = df.groupby('customer service calls')['churn'].mean()

plt.figure(figsize=(10,8))
plt.plot(grouped.index, grouped.values)
plt.xlabel("Customer Service Calls")
plt.ylabel("Churn Rate")
plt.title("Churn Rate by Customer Service Calls")
plt.show()
```



Explanation -For business insight, higher customer service calls might lead to higher churn

```
[9]: # Average minutes
avg_minutes = [
    df['total day minutes'].mean(),
    df['total eve minutes'].mean(),
    df['total night minutes'].mean(),
    df['total intl minutes'].mean()
]

# Average charges
avg_charges = [
    df['total day charge'].mean(),
    df['total eve charge'].mean(),
    df['total night charge'].mean(),
    df['total intl charge'].mean()
]
```

```

labels = ['Day', 'Evening', 'Night', 'International']

x = np.arange(len(labels))
width = 0.35

plt.figure(figsize=(10,6))
bars1 = plt.bar(x - width/2, avg_minutes, width, label='Minutes')
bars2 = plt.bar(x + width/2, avg_charges, width, label='Charges')


plt.xticks(x, labels)
plt.title("Minutes vs Charges Comparison")
plt.ylabel("Average Value")
plt.legend()

# Add values on bars
for bar in bars1:
    plt.text(
        bar.get_x() + bar.get_width()/2,
        bar.get_height(),
        f'{bar.get_height():.1f}',
        ha='center',
        va='bottom'
    )

for bar in bars2:
    plt.text(
        bar.get_x() + bar.get_width()/2,
        bar.get_height(),
        f'{bar.get_height():.2f}',
        ha='center',
        va='bottom'
    )

plt.show()

```



```
[10]: #Clean datafram by converting all categorical data into numeric by use of
      ↪LabelEncoder
      # encoder = LabelEncoder()

      # df_clean= pd.get_dummies(df, drop_first=True)

      # # for col in df.columns:
      # #     df[col] = encoder.fit_transform(df[col])
      # df_clean.info()
```

Separate Target(y) from predictor(X)

```
[11]: X = df.drop('churn', axis=1)
y = df['churn']
```

```
[12]: #selecting important features from the data set using RFE --> Recursive feature
      ↪elimination (target feature selection = 10)

      # lr = LogisticRegression(max_iter=1000)

      # rfe = RFE(lr, n_features_to_select=13)
      # rfe.fit(X, y)

      # selected_features = X.columns[rfe.support_]
```

```
# print(selected_features)
```

Feature Selection

- I began by preparing the data: irrelevant identifiers (e.g. phone number, state and all calls(day, night, eve and intl) were dropped, and categorical fields (“international plan” and “voice mail plan”) were encoded as binary using .getdummies() function.

```
[13]: selected_X = ['account length', 'international plan', 'voice mail plan',
                   'number vmail messages', 'total day minutes', 'total day charge', 'total
                   eve minutes',
                   'total eve charge', 'total night minutes', 'total night charge',
                   'total intl minutes', 'total intl charge', 'customer service calls'
                  ]
X = df[selected_X]
y = df['churn']

cleaned_df = pd.concat([X, y], axis=1)

df = cleaned_df.copy()
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3333 entries, 0 to 3332
Data columns (total 14 columns):
 #   Column           Non-Null Count  Dtype  
 ---  --  
 0   account length    3333 non-null   int64  
 1   international plan 3333 non-null   object  
 2   voice mail plan    3333 non-null   object  
 3   number vmail messages 3333 non-null   int64  
 4   total day minutes   3333 non-null   float64 
 5   total day charge    3333 non-null   float64 
 6   total eve minutes   3333 non-null   float64 
 7   total eve charge    3333 non-null   float64 
 8   total night minutes 3333 non-null   float64 
 9   total night charge   3333 non-null   float64 
 10  total intl minutes   3333 non-null   float64 
 11  total intl charge    3333 non-null   float64 
 12  customer service calls 3333 non-null   int64  
 13  churn              3333 non-null   bool  
dtypes: bool(1), float64(8), int64(3), object(2)
memory usage: 341.9+ KB
```

- Run .describe function to understand statistical information of my data to see if it will need to be normalized

```
[14]: df.describe()
```

```
[14]:      account length  number vmail messages  total day minutes \
count      3333.000000          3333.000000      3333.000000
mean       101.064806          8.099010      179.775098
std        39.822106         13.688365      54.467389
min        1.000000          0.000000      0.000000
25%        74.000000          0.000000     143.700000
50%        101.000000          0.000000     179.400000
75%        127.000000         20.000000     216.400000
max        243.000000         51.000000     350.800000

      total day charge  total eve minutes  total eve charge \
count      3333.000000          3333.000000      3333.000000
mean       30.562307          200.980348      17.083540
std        9.259435          50.713844      4.310668
min        0.000000          0.000000      0.000000
25%        24.430000          166.600000     14.160000
50%        30.500000          201.400000     17.120000
75%        36.790000          235.300000     20.000000
max        59.640000          363.700000     30.910000

      total night minutes  total night charge  total intl minutes \
count      3333.000000          3333.000000      3333.000000
mean       200.872037          9.039325      10.237294
std        50.573847          2.275873      2.791840
min        23.200000          1.040000      0.000000
25%        167.000000          7.520000      8.500000
50%        201.200000          9.050000      10.300000
75%        235.300000          10.590000     12.100000
max        395.000000          17.770000     20.000000

      total intl charge  customer service calls
count      3333.000000          3333.000000
mean       2.764581           1.562856
std        0.753773           1.315491
min        0.000000           0.000000
25%        2.300000           1.000000
50%        2.780000           1.000000
75%        3.270000           2.000000
max        5.400000           9.000000
```

Normalize X using MinMaxScaler

```
[15]: df_clean= pd.get_dummies(cleaned_df, drop_first=True)

df_clean.info()
df = df_clean.copy()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```

RangeIndex: 3333 entries, 0 to 3332
Data columns (total 14 columns):
 #   Column           Non-Null Count  Dtype  
 ---  --  
 0   account length    3333 non-null   int64  
 1   number vmail messages 3333 non-null   int64  
 2   total day minutes  3333 non-null   float64 
 3   total day charge   3333 non-null   float64 
 4   total eve minutes  3333 non-null   float64 
 5   total eve charge   3333 non-null   float64 
 6   total night minutes 3333 non-null   float64 
 7   total night charge  3333 non-null   float64 
 8   total intl minutes 3333 non-null   float64 
 9   total intl charge   3333 non-null   float64 
 10  customer service calls 3333 non-null   int64  
 11  churn              3333 non-null   bool   
 12  international plan_yes 3333 non-null   uint8  
 13  voice mail plan_yes 3333 non-null   uint8  
dtypes: bool(1), float64(8), int64(3), uint8(2)
memory usage: 296.3 KB

```

[16]: `#Separate and assign X and y`

```
X = df.drop('churn', axis=1)
y = df['churn']
```

[17]: `# scaler = MinMaxScaler()`

```
# X_scaled = scaler.fit_transform(X)
```

Train-Test Split

- The data was split into 50% Training and 50% Testing sets.
- Used stratified sampling to ensure the churn distribution remained consistent across both sets.

[18]: `# Split the data into train and test sets.`

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.5,
random_state=254, stratify=y)
```

[19]: `#Normalize using standard scaler`

`#The transformation was then applied to the test set using the training`
`parameters to prevent data leakage.`

```
scaler = StandardScaler()
```

```
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

*Train Models

Four classification models were trained and evaluated: 1. **Logistics Regression**: Linear model used as a baseline, providing interpretability by showing how each feature influences the probability of churn. 2. **Decision Tree Classifier**: Non-linear, rule-based model that can capture complex interactions between features. 3. **Random Forest**: An ensemble of decision trees improving predictive stability and reducing overfitting. 4. **Gradient Boosting**: A boosted ensemble model that iteratively improves weak learners, typically yielding the highest predictive accuracy. All models were trained on the same processed dataset to ensure fair comparison.

1. Logistic Regression Model - Its a classification model used to determine the likelihood of a given data point being associated with one of two categories

```
[20]: lr= LogisticRegression(random_state=254)
lr.fit(X_train_scaled, y_train)

# pipe = Pipeline([
#     ("scaler", StandardScaler()),
#     ("feature_selection", SelectFromModel(LogisticRegression(penalty="l1",
#                                                               solver="liblinear"))),
#     ("model", LogisticRegression())
# ])

# pipe.fit(X_train, y_train)

# y_lrpred = pipe.predict(X_test)
```


```
[20]: LogisticRegression(random_state=254)
```

```
[21]: #Predict model
y_lrpred = lr.predict(X_test_scaled)


#Probability prediction
y_lrprob = lr.predict_proba(X_test_scaled)
```

Plot a heatmap to determine multicollinearity

```
[22]: plt.figure(figsize=(15, 5))
sns.heatmap(df.corr(), annot=True, cmap='Greys', fmt='.2f')
plt.title('Correlation Heatmap')
plt.show()
sns.heatmap(X.corr())
```



[22] : <AxesSubplot:>



- Shows positive and negative drivers of churn

```
[23]: # plt.figure()
# plt.scatter(X, y)    # Actual data points
# plt.plot(X_train_scaled, y_train)  # Logistic curve
# plt.xlabel("Total Day Minutes")
# plt.ylabel("Probability of Churn")
# plt.title("Logistic Regression Curve")
# plt.show()
```

Ridge regression - Its a linear model similar to logistic but with regularization to penalize over estimation

```
[24]: #train
rg = Ridge (alpha =1.0, solver='lsqr') ### penalty strength
rg.fit(X_train_scaled,y_train) #training it on a ridge regression
# (regularization)
y_rgpred = rg.predict(X_test_scaled) # predicting on ridge regression
```

Confirming that Ridge has solved multicollinearity - This will be achieved by comparing their(Logistic and Ridge) coefficients

```
[25]: print(f'logistic coefficient:', lr.coef_)
print()
print(f'Ridge coefficient:', rg.coef_)

logistic coefficient: [[-0.02035763  0.46246376  0.31685152  0.30887201
0.20912403  0.19879638
 0.08285722  0.08346525  0.0773002   0.11140457  0.61280709  0.64507637
-0.81114739]]

Ridge coefficient: [-0.00025419  0.04039593  0.0306566   0.03051825  0.01862427
0.01844627
 0.00754623  0.00750331  0.00888183  0.00943795  0.07004011  0.10106486
-0.06936503]
```

Observation - From the above outcome is evident that ridge regression has solved multicollinearity by shrinking unstable coefficients

2. Decision tree Model - Its a non-linear classifier partitioning features

```
[26]: #Train and fit decision tree model
dt = DecisionTreeClassifier(random_state=254)

dt.fit(X_train_scaled, y_train)
```

```
[26]: DecisionTreeClassifier(random_state=254)
```

```
[27]: #Predict model
y_dtpred = dt.predict(X_test_scaled)
```

```
#Probability prediction  
y_dtprob = dt.predict_proba(X_test_scaled)
```

3. Random Forest Classifier

```
[28]: rfc = RandomForestClassifier(random_state=254)  
  
rfc.fit(X_train_scaled, y_train)
```

```
[28]: RandomForestClassifier(random_state=254)
```

```
[29]: y_rfcpred = rfc.predict(X_test_scaled)  
  
#Probability prediction  
y_rfcpred = rfc.predict_proba(X_test_scaled)
```

4. Gradient boosting classifier


```
[30]: gb = GradientBoostingClassifier(random_state=254)  
  
gb.fit(X_train_scaled, y_train)  
  
y_gbpred = gb.predict(X_test_scaled)  
  
#Probability prediction  
y_gbprob = gb.predict_proba(X_test_scaled)
```

Evaluating our models using the following

- Accuracy: Overall percentage of correct predictions.
- Precision: Percentage of predicted churners who actually churned (reliability of alerts).
- Recall: Percentage of actual churners correctly identified; crucial for reducing revenue loss.
- F1-Score: Harmonic mean of Precision and Recall; balances false positives and false negatives.
- AUC (Area Under ROC Curve): Measures overall ability to distinguish churners from non-churners.

```
[31]: #Confusion matrix  
cm = confusion_matrix(y_test, y_lrpred )  
  
disp = ConfusionMatrixDisplay(confusion_matrix=cm)  
disp.plot()
```

```
[31]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x23221f96ee0>
```



[32]: #Accuracy Score

```
accuracylr = accuracy_score(y_test,y_lrpred)
accuracydt = accuracy_score(y_test,y_dtpred)
accuracyrfc = accuracy_score(y_test,y_rfcpred)
accuracyrgb = accuracy_score(y_test,y_gbpred)

print(f'Accuracy score of the logistic model fitted:{accuracylr:.3f} {accuracylr*100:.1f}%')
print()
print(f'Accuracy score of the decision tree fitted:{accuracydt:.3f} {accuracydt*100:.1f}%')
print()
print(f'Accuracy score of the random forest fitted:{accuracyrfc:.3f} {accuracyrfc*100:.1f}%')
print()
print(f'Accuracy score of the gradient boosting fitted:{accuracyrfc:.3f} {accuracyrfc*100:.1f}%')
```

Accuracy score of the logistic model fitted:0.855 (85.5%)

Accuracy score of the decision tree fitted:0.888 (88.8%)

Accuracy score of the random forest fitted:0.932 (93.2%)

Accuracy score of the gradient boosting fitted:0.932 (93.2%)

2 Interpretation:

- Random Forest and Gradient Boosting have the highest overall accuracy. However, due to class imbalance, accuracy alone does not fully reflect the ability to detect churners.

[33]: #Compare clasification report for Logistic model and Decision tree --> to
↳easily compare the scores

```
print(f'Classification report model for Logistic regression')
print(classification_report(y_test, y_lrpred))
print()
print(f'Classification report model for Decision tree')
print(classification_report(y_test, y_dtpred))
print()
print(f'Classification report model for Random forest')
print(classification_report(y_test, y_rfcpred))
print()
print(f'Classification report model for Gradient boosting')
print(classification_report(y_test, y_gbpred))
```

Classification report model for Logistic regression

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

False	0.87	0.97	0.92	1425
True	0.51	0.18	0.26	242
accuracy			0.86	1667
macro avg	0.69	0.57	0.59	1667
weighted avg	0.82	0.86	0.82	1667

Classification report model for Decision tree

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

False	0.93	0.94	0.93	1425
True	0.62	0.60	0.61	242
accuracy			0.89	1667
macro avg	0.77	0.77	0.77	1667
weighted avg	0.89	0.89	0.89	1667

Classification report model for Random forest

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

False	0.93	0.99	0.96	1425
True	0.93	0.58	0.71	242
accuracy			0.93	1667
macro avg	0.93	0.79	0.84	1667
weighted avg	0.93	0.93	0.93	1667

Classification report model for Gradient boosting				
	precision	recall	f1-score	support
False	0.94	0.98	0.96	1425
True	0.86	0.62	0.72	242
accuracy			0.93	1667
macro avg	0.90	0.80	0.84	1667
weighted avg	0.93	0.93	0.92	1667

Evaluating Decision tree

The following will be used for model 1. accuracy score 2. Roc curve 3. Auc curve

```
[34]: print(f'Precision score for Logistic Regression Model --->',  
      ↪precision_score(y_test, y_lrpred))  
print(f'Recall score for Logistic Regression Model --->', recall_score(y_test,  
      ↪y_lrpred))  
print(f'F1 score for Logistic Regression Model --->', f1_score(y_test,  
      ↪y_lrpred))  
print()  
print(f'Precision score for Decision tree Model --->', precision_score(y_test,  
      ↪y_dtpred))  
print(f'Recall score for Decision tree Model --->', recall_score(y_test,  
      ↪y_dtpred))  
print(f'F1 score for Decision tree Model --->', f1_score(y_test, y_dtpred))  
print()  
print(f'Precision score for Random forest Model --->', precision_score(y_test,  
      ↪y_rfcpred))  
print(f'Recall score for Random forest Model --->', recall_score(y_test,  
      ↪y_rfcpred))  
print(f'F1 score for Random forest Model --->', f1_score(y_test, y_rfcpred))  
print()  
print(f'Precision score for Gradient boosting Model --->',  
      ↪precision_score(y_test, y_gbpred))  
print(f'Recall score for Gradient boosting --->', recall_score(y_test,  
      ↪y_gbpred))  
print(f'F1 score for Gradient boosting Model --->', f1_score(y_test, y_gbpred))
```

```
Precision score for Logistic Regression Model ---> 0.5058823529411764
Recall score for Logistic Regression Model ---> 0.17768595041322313
F1 score for Logistic Regression Model ---> 0.26299694189602446
```

```
Precision score for Decision tree Model ---> 0.6180257510729614
Recall score for Decision tree Model ---> 0.5950413223140496
F1 score for Decision tree Model ---> 0.6063157894736843
```

```
Precision score for Random forest Model ---> 0.9271523178807947
Recall score for Random forest Model ---> 0.5785123966942148
F1 score for Random forest Model ---> 0.7124681933842238
```

```
Precision score for Gradient boosting Model ---> 0.861271676300578
Recall score for Gradient boosting ---> 0.6157024793388429
F1 score for Gradient boosting Model ---> 0.7180722891566265
```

3 Interpretation:

- **Logistic Regression:** Misses most churners (Recall = 17.8%), making it unsuitable for proactive retention despite decent accuracy.
- **Decision Tree:** Balanced performance (Recall = 59.5%, F1 = 60.6%), interpretable rules for business use.
- **Random Forest:** Extremely precise (92.7%) and strong F1-score (71.2%), but moderate recall means some churners are still missed.
- **Gradient Boosting:** Best overall balance (Recall = 61.6%, F1 = 71.8%), high AUC (0.913), and strong predictive performance.

```
[35]: import matplotlib.pyplot as plt
import numpy as np

# Define metrics and their values
metrics = {
    "Accuracy": 93.2,      # in percentage
    "Precision": 86.1,
    "Recall": 61.6,
    "F1-Score": 71.8,
    "AUC": 91.3
}

# Function to draw a simple gauge
def draw_gauge(ax, value, label):
    # Draw background semicircle
    theta = np.linspace(-np.pi/2, np.pi/2, 100)
    ax.plot(np.cos(theta), np.sin(theta), color='lightgray', linewidth=20, solid_capstyle='round')

    # Draw value arc
```

```

theta_value = np.linspace(-np.pi/2, -np.pi/2 + (value/100)*np.pi, 100)
ax.plot(np.cos(theta_value), np.sin(theta_value), color="#1f77b4", linewidth=20, solid_capstyle='round')

# Draw needle
needle_angle = -np.pi/2 + (value/100)*np.pi
ax.plot([0, np.cos(needle_angle)*0.8], [0, np.sin(needle_angle)*0.8], color='red', linewidth=3)

# Add text
ax.text(0, -0.2, f"{label}\n{value:.1f}%", ha='center', va='center', fontsize=12, weight='bold')


# Hide axes
ax.axis('off')
ax.set_aspect('equal')

# Create figure with subplots for each metric
fig, axes = plt.subplots(1, len(metrics), figsize=(18,4))

for ax, (metric, value) in zip(axes, metrics.items()):
    draw_gauge(ax, value, metric)

plt.suptitle("Churn Prediction Metrics Overview", fontsize=16, weight='bold')
plt.tight_layout()
plt.show()

```



```
[36]: # Models and recall values
models = ["Logistic Regression", "Decision Tree", "Random Forest", "GradientBoosting"]
recall_values = [17.77, 59.50, 57.85, 61.57] # in percentage

# Colors: highlight Gradient Boosting
colors = ['skyblue', 'skyblue', 'skyblue', 'orange']

# Create bar chart
```

```


plt.figure(figsize=(8,5))
bars = plt.bar(models, recall_values, color=colors)

# Annotate values on top of bars
for bar in bars:
    yval = bar.get_height()
    plt.text(bar.get_x() + bar.get_width()/2, yval + 1, f"{yval:.1f}%", ha='center', va='bottom', fontsize=10)

# Chart title and labels
plt.title("Recall Comparison by Model", fontsize=14, weight='bold')
plt.ylabel("Recall (%)")
plt.ylim(0, 70)

# Show chart
plt.tight_layout()
plt.show()

```



[37]: #Accuracy score

```

accuracy_dt = accuracy_score(y_test, y_dtpred)
accuracy_dt

```

[37]: 0.8878224355128974

```
[38]: #To determin AUC and ROC curve
y_lrprob = lr.predict_proba(X_test_scaled)[:,1]
y_dtprob = dt.predict_proba(X_test_scaled)[:,1]
y_rfcprob = rfc.predict_proba(X_test_scaled)[:,1]
y_gbprob = gb.predict_proba(X_test_scaled)[:,1]

lr_fpr, lr_tpr, i = roc_curve(y_test, y_lrprob)
dt_fpr, dt_tpr, i = roc_curve(y_test, y_dtprob)
rfc_fpr, rfc_tpr, i = roc_curve(y_test, y_rfcprob)
gb_fpr, gb_tpr, i = roc_curve(y_test, y_gbprob)

lr_auc = auc(lr_fpr, lr_tpr)
dt_auc = auc(dt_fpr, dt_tpr)
rfc_auc = auc(rfc_fpr, rfc_tpr)
gb_auc = auc(gb_fpr, gb_tpr)


print(f'Logistic Regression AUC:{lr_auc},')
print('Decision tree AUC:',dt_auc)
print('Random random forest AUC:',rfc_auc)
print('gradient boosting AUC:',gb_auc)
```

```
Logistic Regression AUC: {0.8301348412353197}
Decision tree AUC: {0.7662925909815862}
Random random forest AUC: {0.9060939538929971}
gradient boosting AUC: {0.9128925619834711}
```

```
[39]: plt.figure(figsize=(8, 6))
plt.plot(lr_fpr, lr_tpr, label=f'LogisticRegression (AUC = {lr_auc:2f})')
plt.plot(dt_fpr, dt_tpr, label=f'Decision Tree AUC = {dt_auc:2f}')
plt.plot(rfc_fpr, rfc_tpr, label=f'Random forest AUC = {rfc_auc:2f}')
plt.plot(gb_fpr, gb_tpr, label=f'Gradient boosting AUC = {gb_auc:2f}')
plt.plot([0,1], [0, 1], linestyle='--')

plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic (ROC) Curve')

plt.legend()
plt.show()
```



3.1 ROC Curve Visualization

4 Interpretation

- Gradient Boosting demonstrates the strongest ability to separate churners from non-churners, followed closely by Random Forest. Decision Tree and Logistic Regression are less effective at distinguishing classes.

4.1 Model types and their increasing complexity

```
[40]: import matplotlib.pyplot as plt
from matplotlib.patches import FancyArrow

# Define models and complexity levels
models = ["Logistic Regression", "Decision Tree", "Random Forest", "GradientBoosting"]
complexity = [1, 2, 3, 4] # increasing complexity

# Colors for each model
colors = ["#1f77b4", "#ff7f0e", "#2ca02c", "#d62728"]
```

```

# Create figure
fig, ax = plt.subplots(figsize=(10, 3))
ax.set_xlim(0, 5)
ax.set_ylim(0, 2)
ax.axis("off")

# Plot models as circles with arrows showing increasing complexity
for i, model in enumerate(models):
    # Circle for model
    circle = plt.Circle((complexity[i], 1), 0.2, color=colors[i], alpha=0.8)
    ax.add_patch(circle)

    # Label model
    ax.text(complexity[i], 1, model, color='white', fontsize=10, ha='center', va='center', weight='bold')

    # Draw arrow to next model (except the last)
    if i < len(models)-1:
        arrow = FancyArrow(complexity[i]+0.2, 1, 0.6, 0, width=0.05, length_includes_head=True, color='gray', alpha=0.6)
        ax.add_patch(arrow)

# Title
plt.title("Model Types and Increasing Complexity", fontsize=14)
plt.show()

```

Model Types and Increasing Complexity



5 Limitations

- Class Imbalance:** Fewer churners than retained customers may cause models to under-detect minority cases. Some churners could still be missed despite using ensemble methods.
- Threshold Dependence:** Metrics are based on a standard probability threshold (0.5). Adjusting thresholds could improve recall at the expense of precision.

3. **Feature Limitations:** Dataset contains limited features. External factors such as competitor pricing, promotions, or customer sentiment were not included and may impact churn behavior.
4. **Overfitting Risk:** Complex models like Random Forest and Gradient Boosting may overfit to training data if not monitored carefully, reducing generalization on new customer data.
5. **Interpretability:** Ensemble models (Random Forest, Gradient Boosting) provide less transparency compared to Decision Trees or Logistic Regression, which may complicate explaining decisions to non-technical stakeholders.

6 Conclusion

Key Insights - Logistic Regression: Poor recall, not suitable for churn prevention. - **Decision Tree:** Interpretable, moderate performance; good for explaining patterns to business teams. - **Random Forest:** Highly precise; excellent for targeted interventions. - **Gradient Boosting:** Best balance of recall, precision, F1-score, and AUC; most effective for proactive retention.

7 Recommendations

The following are my recommendation to the firm as per the metric evaluation tests done

1. Deploy Gradient Boosting for production due to superior overall performance.
2. Use predicted probabilities to rank high-risk customers for targeted retention campaigns.
3. Integrate predictions into CRM systems for real-time alerts to account managers.
4. Adjust thresholds to increase recall if business priorities favor identifying more churners.
5. Future Work: Explore additional features, resampling methods, or alternative ensemble techniques to further improve detection of minority churn cases.